

Practice Lab: Linear Regression

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import copy
import math
```

```
In [2]: #x_train is the population of a city
#y_train is the profit of a restaurant in that city. A negative value for profit indic
# Both X_train and y_train are numpy arrays.

x_train, y_train = load_data(...)
```

```
In [3]: # print x_train
print("Type of x_train:", type(x_train))
print("First five elements of x_train are:\n", x_train[:5])

Type of x_train: <class 'numpy.ndarray'>
First five elements of x_train are:
[6.1101 5.5277 8.5186 7.0032 5.8598]
```

```
In [4]: # print y_train
print("Type of y_train:", type(y_train))
print("First five elements of y_train are:\n", y_train[:5])


Type of y_train: <class 'numpy.ndarray'>
First five elements of y_train are:
[17.592  9.1302 13.662 11.854  6.8233]
```

```
In [5]: print ('The shape of x_train is:', x_train.shape)
print ('The shape of y_train is: ', y_train.shape)
print ('Number of training examples (m):', len(x_train))

The shape of x_train is: (97,)
The shape of y_train is: (97,)
Number of training examples (m): 97
```

```
In [6]: # Create a scatter plot of the data. To change the markers to red "x",
# we used the 'marker' and 'c' parameters
plt.scatter(x_train, y_train, marker='x', c='r')

# Set the title
plt.title("Profits vs. Population per city")
# Set the y-axis label
plt.ylabel('Profit in $10,000')
# Set the x-axis label
plt.xlabel('Population of City in 10,000s')
plt.show()
```



Compute Cost function

```
In [7]: # compute cost function fwb(x) = wx + b

def compute_cost(x, y, w, b):
    """
    Computes the cost function for linear regression.

    Args:
        x (ndarray): Shape (m,) Input to the model (Population of cities)
        y (ndarray): Shape (m,) Label (Actual profits for the cities)
        w, b (scalar): Parameters of the model

    Returns
        total_cost (float): The cost of using w,b as the parameters for linear regression
                           to fit the data points in x and y
    """
    # number of training examples
    m = x.shape[0]
    total_cost = 0

    cum_cost = 0
    for i in range(m):
        f_wb = np.dot(x[i], w) + b
        cost = (f_wb - y[i])**2
        cum_cost = cum_cost + cost
    total_cost = cum_cost / (2*m)

    return total_cost
```

Gradient descent

```
In [9]: # compute_gradient

def compute_gradient(x, y, w, b):
    """
    Computes the gradient for linear regression
    Args:
        x (ndarray): Shape (m,) Input to the model (Population of cities)
        y (ndarray): Shape (m,) Label (Actual profits for the cities)
```

```

    w, b (scalar): Parameters of the model
Returns
    dj_dw (scalar): The gradient of the cost w.r.t. the parameters w
    dj_db (scalar): The gradient of the cost w.r.t. the parameter b
"""

# Number of training examples
m = x.shape[0]

# returning dj_dw, dj_db
dj_dw = 0
dj_db = 0
for i in range(m):
    f_wb=w*x[i]+b
    dj_dw_i= (f_wb-y[i])*x[i]
    dj_db_i= f_wb-y[i]
    dj_dw = dj_dw + dj_dw_i
    dj_db = dj_db + dj_db_i
dj_dw = dj_dw/m
dj_db = dj_db/m
return dj_dw, dj_db

```

```

In [11]: # Compute and display cost and gradient with random w
test_w = 0.2
test_b = 0.2
tmp_dj_dw, tmp_dj_db = compute_gradient(x_train, y_train, test_w, test_b)

print('Gradient at test w, b:', tmp_dj_dw, tmp_dj_db)

```

Gradient at test w, b: -47.41610118114435 -4.007175051546391

Learning parameters using batch gradient descent

```

In [12]: def gradient_descent(x, y, w_in, b_in, cost_function, gradient_function, alpha, num_it
"""
    Performs batch gradient descent to learn theta. Updates theta by taking
    num_iters gradient steps with learning rate alpha

    Args:
        x : (ndarray): Shape (m,)
        y : (ndarray): Shape (m,)
        w_in, b_in : (scalar) Initial values of parameters of the model
        cost_function: function to compute cost
        gradient_function: function to compute the gradient
        alpha : (float) Learning rate
        num_iters : (int) number of iterations to run gradient descent
    Returns
        w : (ndarray): Shape (1,) Updated values of parameters of the model after
            running gradient descent
        b : (scalar) Updated value of parameter of the model after
            running gradient descent
"""

# number of training examples
m = len(x)

# An array to store cost J and w's at each iteration – primarily for graphing later
J_history = []
w_history = []

```

```

w = copy.deepcopy(w_in) #avoid modifying global w within function
b = b_in

for i in range(num_iters):

    # Calculate the gradient and update the parameters
    dj_dw, dj_db = gradient_function(x, y, w, b )

    # Update Parameters using w, b, alpha and gradient
    w = w - alpha * dj_dw
    b = b - alpha * dj_db

    # Save cost J at each iteration
    if i<100000:      # prevent resource exhaustion
        cost = cost_function(x, y, w, b)
        J_history.append(cost)

    # Print cost every at intervals 10 times or as many iterations if < 10
    if i% math.ceil(num_iters/10) == 0:
        w_history.append(w)
        print(f"Iteration {i:4}: Cost {float(J_history[-1]):8.2f}  ")

return w, b, J_history, w_history #return w and J,w history for graphing

```

```

In [13]: # initialize fitting parameters
initial_w = 0.
initial_b = 0.

# some gradient descent settings
iterations = 1500
alpha = 0.01

w,b,_,_ = gradient_descent(x_train ,y_train, initial_w, initial_b,
                           compute_cost, compute_gradient, alpha, iterations)
print("w,b found by gradient descent:", w, b)

```

```

Iteration    0: Cost      6.74
Iteration   150: Cost      5.31
Iteration   300: Cost      4.96
Iteration   450: Cost      4.76
Iteration   600: Cost      4.64
Iteration   750: Cost      4.57
Iteration   900: Cost      4.53
Iteration  1050: Cost      4.51
Iteration  1200: Cost      4.50
Iteration  1350: Cost      4.49
w,b found by gradient descent: 1.166362350335582 -3.63029143940436

```

```

In [14]: # use the generated w,b to calculate the predictions
m = x_train.shape[0]
predicted = np.zeros(m)

for i in range(m):
    predicted[i] = w * x_train[i] + b

```

```

In [15]: # Plot the linear fit
plt.plot(x_train, predicted, c = "b")


# Create a scatter plot of the data.

```

```
plt.scatter(x_train, y_train, marker='x', c='r')

# Set the title
plt.title("Profits vs. Population per city")
# Set the y-axis label
plt.ylabel('Profit in $10,000')
# Set the x-axis label
plt.xlabel('Population of City in 10,000s')
```

Out[15]: Text(0.5, 0, 'Population of City in 10,000s')



```
In [16]: # predict the profit in the areas of 35,000 and 70,000 people
predict1 = 3.5 * w + b
print('For population = 35,000, we predict a profit of $%.2f' % (predict1*10000))

predict2 = 7.0 * w + b
print('For population = 70,000, we predict a profit of $%.2f' % (predict2*10000))
```

For population = 35,000, we predict a profit of \$4519.77
 For population = 70,000, we predict a profit of \$45342.45