

# Practice lab: Logistic Regression

In this exercise, we will implement logistic regression and apply it to two different datasets.

```
In [13]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import copy
import math
```

```
In [14]: # Load dataset
"""
X_train contains exam scores on two exams for a student
y_train is the admission decision
y_train = 1 if the student was admitted
y_train = 0 if the student was not admitted
Both X_train and y_train are numpy arrays.
"""

df = pd.read_csv(r'...')
X_train, y_train = df.to_numpy()
```

```
In [15]: print("First five elements in X_train are:\n", X_train[:5])
print("Type of X_train:", type(X_train))
```

```
First five elements in X_train are:
[[34.62365962 78.02469282]
 [30.28671077 43.89499752]
 [35.84740877 72.90219803]
 [60.18259939 86.3085521 ]
 [79.03273605 75.34437644]]
Type of X_train: <class 'numpy.ndarray'>
```

```
In [16]: print("First five elements in y_train are:\n", y_train[:5])
print("Type of y_train:", type(y_train))
```

```
First five elements in y_train are:
[0. 0. 0. 1. 1.]
Type of y_train: <class 'numpy.ndarray'>
```


```
In [17]: print ('The shape of X_train is: ' + str(X_train.shape))
print ('The shape of y_train is: ' + str(y_train.shape))
print ('We have m = %d training examples' % (len(y_train)))
```

```
The shape of X_train is: (100, 2)
The shape of y_train is: (100,)
We have m = 100 training examples
```

```
In [18]: # Plot examples
plot_data(X_train, y_train[:,], pos_label="Admitted", neg_label="Not admitted")

# Set the y-axis label
plt.ylabel('Exam 2 score')
# Set the x-axis label
plt.xlabel('Exam 1 score')
```

```
plt.legend(loc="upper right")
plt.show()
```



In [22]: *# implement sigmoid function*

```
def sigmoid(z):
    """
    Compute the sigmoid of z

    Args:
        z (ndarray): A scalar, numpy array of any size.

    Returns:
        g (ndarray): sigmoid(z), with the same shape as z

    """
    g = 1/(1+np.exp(-z))
    return g
```

In [33]: *# implement compute\_cost*

```
def compute_cost(X, y, w, b, lambda_= 1):
    """
    Computes the cost over all examples

    Args:
        X : (ndarray Shape (m,n)) data, m examples by n features
        y : (array_like Shape (m,)) target value
        w : (array_like Shape (n,)) Values of parameters of the model
        b : scalar Values of bias parameter of the model
        lambda_: unused placeholder

    Returns:
        total_cost: (scalar)        cost
    """
    m, n = X.shape

    loss_sum = 0
    for i in range(m):
        z_i = 0
        for j in range(n):
            # Add the corresponding term to z_wb
            z_ij = np.dot(X[i,j],w[j])
            z_i +=z_ij
        # Add the bias term to z_wb
```

```

        z_i = z_i + b
        # Add the sigmoid function to z_wb
        f_wb_i = sigmoid(z_i)
        loss = -y[i]*np.log(f_wb_i) - (1-y[i])*np.log(1-f_wb_i)
        loss_sum += loss
    total_cost = loss_sum/m

    return total_cost

```

In [34]: `m, n = X_train.shape`

```

# Compute and display cost with w initialized to zeroes
initial_w = np.zeros(n)
initial_b = 0.
cost = compute_cost(X_train, y_train, initial_w, initial_b)
print('Cost at initial w (zeros): {:.3f}'.format(cost))

```

Cost at initial w (zeros): 0.693

In [48]:

```

# implement gradient descent function
def compute_gradient(X, y, w, b, lambda_=None):
    """
    Computes the gradient for logistic regression

    Args:
        X : (ndarray Shape (m,n)) variable such as house size
        y : (array_like Shape (m,1)) actual value
        w : (array_like Shape (n,1)) values of parameters of the model
        b : (scalar) value of parameter of the model
        lambda_: unused placeholder.

    Returns
        dj_dw: (array_like Shape (n,1)) The gradient of the cost w.r.t. the parameters w
        dj_db: (scalar) The gradient of the cost w.r.t. the parameter b.
    """

    m, n = X.shape
    dj_dw = np.zeros(w.shape)
    dj_db = 0.

    for i in range(m):
        z_wb = 0
        for j in range(n):
            z_wb += np.dot(X[i,j],w[j])
        z_wb += b
        f_wb = sigmoid(z_wb)

        dj_db_i = f_wb - y[i]
        dj_db += dj_db_i

        for j in range(n):
            dj_dw[j] += (f_wb - y[i])*X[i,j]

    dj_dw = (1/m)*dj_dw
    dj_db = (1/m)*dj_db

    return dj_db, dj_dw

```

```
In [49]: # Compute and display gradient with w initialized to zeroes
initial_w = np.zeros(n)
initial_b = 0.

dj_db, dj_dw = compute_gradient(X_train, y_train, initial_w, initial_b)
print(f'dj_db at initial w (zeros):{dj_db}' )
print(f'dj_dw at initial w (zeros):{dj_dw.tolist()}' )

dj_db at initial w (zeros):-0.1
dj_dw at initial w (zeros):[-12.00921658929115, -11.262842205513591]
```

## Learning parameters using gradient descent

```
In [51]: def gradient_descent(X, y, w_in, b_in, cost_function, gradient_function, alpha, num_it
"""
    Performs batch gradient descent to learn theta. Updates theta by taking
    num_iters gradient steps with learning rate alpha

    Args:
        X : (array_like Shape (m, n))
        y : (array_like Shape (m,))
        w_in : (array_like Shape (n,)) Initial values of parameters of the model
        b_in : (scalar) Initial value of parameter of the model
        cost_function: function to compute cost
        alpha : (float) Learning rate
        num_iters : (int) number of iterations to run gradient descent
        lambda_ (scalar, float) regularization constant

    Returns:
        w : (array_like Shape (n,)) Updated values of parameters of the model after
            running gradient descent
        b : (scalar) Updated value of parameter of the model after
            running gradient descent
    """

    # number of training examples
    m = len(X)

    # An array to store cost J and w's at each iteration for graphing
    J_history = []
    w_history = []

    for i in range(num_iters):

        # Calculate the gradient and update the parameters
        dj_db, dj_dw = gradient_function(X, y, w_in, b_in, lambda_)

        # Update Parameters using w, b, alpha and gradient
        w_in = w_in - alpha * dj_dw
        b_in = b_in - alpha * dj_db

        # Save cost J at each iteration
        if i < 100000: # prevent resource exhaustion
            cost = cost_function(X, y, w_in, b_in, lambda_)
            J_history.append(cost)

        # Print cost every at intervals 10 times or as many iterations if < 10
        if i % math.ceil(num_iters/10) == 0 or i == (num_iters-1):
```

```

w_history.append(w_in)
print(f"Iteration {i:4}: Cost {float(J_history[-1]):8.2f}  ")

return w_in, b_in, J_history, w_history #return w and J,w history for graphing

```

```

In [52]: # run gradient descent to learn the parameters
np.random.seed(1)
initial_w = 0.01 * (np.random.rand(2).reshape(-1,1) - 0.5)
initial_b = -8

# Some gradient descent settings
iterations = 10000
alpha = 0.001

w,b, J_history,_ = gradient_descent(X_train ,y_train, initial_w, initial_b,
                                   compute_cost, compute_gradient, alpha, iterations,

```

```

Iteration    0: Cost      1.01
Iteration 1000: Cost      0.31
Iteration 2000: Cost      0.30
Iteration 3000: Cost      0.30
Iteration 4000: Cost      0.30
Iteration 5000: Cost      0.30
Iteration 6000: Cost      0.30
Iteration 7000: Cost      0.30
Iteration 8000: Cost      0.30
Iteration 9000: Cost      0.30
Iteration 9999: Cost      0.30


```

## Plotting the decision boundary

```

In [53]: plot_decision_boundary(w, b, X_train, y_train)

```



## Evaluating logistic regression

```

In [56]: # implement predict function to produce 1 or 0 predictions given a dataset and w,b

def predict(X, w, b):
    """
    Predict whether the label is 0 or 1 using learned logistic

```

regression parameters w

Args:

X : (ndarray Shape (m, n))

w : (array\_like Shape (n,))

Parameters of the model

b : (scalar, float)

Parameter of the model

Returns:

p: (ndarray (m,1))

The predictions for X using a threshold at 0.5

"""

*# number of training examples*

m, n = X.shape

p = np.zeros(m)

*# Loop over each example*

for i in range(m):

z\_wb = 0

*# Loop over each feature*

for j in range(n):

*# Add the corresponding term to z\_wb*

z\_wb += np.dot(X[i,j],w[j])

*# Add bias term*

z\_wb += b

*# Calculate the prediction for this example*

f\_wb = sigmoid(z\_wb)

*# Apply the threshold (here it is 0.5)*

p[i] = 1 if f\_wb>=0.5 else 0

return p

In [58]: *#Compute accuracy on training set*

p = predict(X\_train, w,b)

print('Train Accuracy: %f'%(np.mean(p == y\_train) \* 100))

Train Accuracy: 92.000000

## Regularized Logistic Regression

In this part of the exercise, we implement regularized logistic regression to predict whether microchips from a fabrication plant passes quality assurance (QA). During QA, each microchip goes through various tests to ensure it is functioning correctly.

In [59]: *# Load dataset*

df = pd.read\_csv(r'...')

X\_train, y\_train = df.to\_numpy()

In [60]: *# print X\_train*

print("X\_train:", X\_train[:5])

print("Type of X\_train:",type(X\_train))

*# print y\_train*

print("y\_train:", y\_train[:5])

print("Type of y\_train:",type(y\_train))


```
X_train: [[ 0.051267  0.69956 ]
 [-0.092742  0.68494 ]
 [-0.21371   0.69225 ]
 [-0.375     0.50219 ]
 [-0.51325   0.46564 ]]
Type of X_train: <class 'numpy.ndarray'>
y_train: [1. 1. 1. 1. 1.]
Type of y_train: <class 'numpy.ndarray'>
```

```
In [61]: print ('The shape of X_train is: ' + str(X_train.shape))
print ('The shape of y_train is: ' + str(y_train.shape))
print ('We have m = %d training examples' % (len(y_train)))
```

```
The shape of X_train is: (118, 2)
The shape of y_train is: (118,)
We have m = 118 training examples
```

```
In [62]: # Plot examples
plot_data(X_train, y_train[:,], pos_label="Accepted", neg_label="Rejected")

# Set the y-axis label
plt.ylabel('Microchip Test 2')
# Set the x-axis label
plt.xlabel('Microchip Test 1')
plt.legend(loc="upper right")
plt.show()
```



## Feature mapping

feature mapping allows us to build a more expressive classifier

```
In [63]: print("Original shape of data:", X_train.shape)

mapped_X = map_feature(X_train[:, 0], X_train[:, 1])
print("Shape after feature mapping:", mapped_X.shape)
```

```
Original shape of data: (118, 2)
Shape after feature mapping: (118, 27)
```

```
In [64]: print("X_train[0]:", X_train[0])
print("mapped X_train[0]:", mapped_X[0])
```

```
X_train[0]: [0.051267 0.69956 ]
mapped X_train[0]: [5.12670000e-02 6.99560000e-01 2.62830529e-03 3.58643425e-02
4.89384194e-01 1.34745327e-04 1.83865725e-03 2.50892595e-02
3.42353606e-01 6.90798869e-06 9.42624411e-05 1.28625106e-03
1.75514423e-02 2.39496889e-01 3.54151856e-07 4.83255257e-06
6.59422333e-05 8.99809795e-04 1.22782870e-02 1.67542444e-01
1.81563032e-08 2.47750473e-07 3.38066048e-06 4.61305487e-05
6.29470940e-04 8.58939846e-03 1.17205992e-01]
```

## 3.4 Cost function for regularized logistic regression

feature mapping is more susceptible to overfitting, we will implement regularized logistic regression to fit the data and regularize them to prevent overfitting.

```
In [65]: # implement cost function for regularized logistic regression
def compute_cost_reg(X, y, w, b, lambda_ = 1):
    """
    Computes the cost over all examples
    Args:
        X : (array_like Shape (m,n)) data, m examples by n features
        y : (array_like Shape (m,)) target value
        w : (array_like Shape (n,)) Values of parameters of the model
        b : (array_like Shape (n,)) Values of bias parameter of the model
        lambda_ : (scalar, float) Controls amount of regularization
    Returns:
        total_cost: (scalar) cost
    """

    m, n = X.shape

    # Calls the compute_cost function that you implemented above
    cost_without_reg = compute_cost(X, y, w, b)

    reg_cost = 0.

    for j in range(n):
        reg_cost += w[j]**2

    # Add the regularization cost to get the total cost
    total_cost = cost_without_reg + (lambda_/(2 * m)) * reg_cost

    return total_cost
```

## Gradient for regularized logistic regression

```
In [ ]: # implement gradient descent for regularized logistic regression
def compute_gradient_reg(X, y, w, b, lambda_ = 1):
    """
    Computes the gradient for linear regression

    Args:
        X : (ndarray Shape (m,n)) variable such as house size
        y : (ndarray Shape (m,)) actual value
        w : (ndarray Shape (n,)) values of parameters of the model
        b : (scalar) value of parameter of the model
        lambda_ : (scalar,float) regularization constant
    Returns
```



```

    dj_db: (scalar)           The gradient of the cost w.r.t. the parameter b.
    dj_dw: (ndarray Shape (n,)) The gradient of the cost w.r.t. the parameters w.

    """
    m, n = X.shape

    dj_db, dj_dw = compute_gradient(X, y, w, b)

    for j in range(n):
        dj_dw_j_reg = w[j] * lambda_ / m
        dj_dw[j] = dj_dw[j] + dj_dw_j_reg
    return dj_db, dj_dw

```

```

In [83]: ## Learning parameters using gradient descent
# Initialize fitting parameters
np.random.seed(1)
initial_w = np.random.rand(X_mapped.shape[1]) - 0.5
initial_b = 1.

# Set regularization parameter lambda_ (you can try varying this)
lambda_ = 0.01

# Some gradient descent settings
iterations = 10000
alpha = 0.01

w, b, J_history, _ = gradient_descent(X_mapped, y_train, initial_w, initial_b,
                                     compute_cost_reg, compute_gradient_reg,
                                     alpha, iterations, lambda_)

```

```


Iteration   0: Cost    0.72
Iteration 1000: Cost    0.59
Iteration 2000: Cost    0.56
Iteration 3000: Cost    0.53
Iteration 4000: Cost    0.51
Iteration 5000: Cost    0.50
Iteration 6000: Cost    0.48
Iteration 7000: Cost    0.47
Iteration 8000: Cost    0.46
Iteration 9000: Cost    0.45
Iteration 9999: Cost    0.45

```

```

In [84]: plot_decision_boundary(w, b, X_mapped, y_train)

```



```
In [85]: #Compute accuracy on the training set
p = predict(X_mapped, w, b)

print('Train Accuracy: %f'%(np.mean(p == y_train) * 100))

Train Accuracy: 82.203390
```